

XGBoost_Salary_Classifier

December 26, 2020

```
[1]: import pandas as pd
import numpy as np
import scipy.stats as st
from sklearn.ensemble import GradientBoostingClassifier
```

```
[2]: train_set=pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/
    ↳adult/adult.data',header=None)
```

```
[91]: test_set=pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/
    ↳adult/adult.test',skiprows=1,header=None)
```

```
[92]: col_labels = ['age', 'workclass', 'fnlwgt', 'education',
    ↳'education_num', 'marital_status', 'occupation', 'relationship', 'race',
    ↳'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
    ↳'wage_class']
```

```
[93]: train_set.columns=col_labels
test_set.columns=col_labels
```

```
[94]: #Getting a glimpse of train data
train_set.head()
```

```
[94]:
```

	age	workclass	fnlwgt	education	education_num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital_status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	1	
1	Married-civ-spouse	Exec-managerial	Husband	White	1	
2	Divorced	Handlers-cleaners	Not-in-family	White	1	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	1	
4	Married-civ-spouse	Prof-specialty	Wife	Black	0	

	capital_gain	capital_loss	hours_per_week	native_country	wage_class
--	--------------	--------------	----------------	----------------	------------

0	2174	0	40	1	1
1	0	0	13	1	1
2	0	0	40	1	1
3	0	0	40	1	1
4	0	0	40	0	1

```
[95]: #Getting a Glimpse of test data
test_set.head()
```

```
[95]:   age  workclass  fnlwgt   education  education_num  marital_status \
0   25   Private  226802      11th           7      Never-married
1   38   Private   89814     HS-grad           9  Married-civ-spouse
2   28  Local-gov  336951  Assoc-acdm          12  Married-civ-spouse
3   44   Private  160323  Some-college          10  Married-civ-spouse
4   18      ?    103497  Some-college          10      Never-married

      occupation  relationship   race   sex  capital_gain \
0  Machine-op-inspct   Own-child  Black  Male           0
1   Farming-fishing   Husband  White  Male           0
2   Protective-serv   Husband  White  Male           0
3  Machine-op-inspct   Husband  Black  Male       7688
4      ?      Own-child  White  Female           0

      capital_loss  hours_per_week  native_country  wage_class
0           0           40  United-States  <=50K.
1           0           50  United-States  <=50K.
2           0           40  United-States  >50K.
3           0           40  United-States  >50K.
4           0           30  United-States  <=50K.
```

```
[8]: #Getting shape details of both train and test data
print('Train Dataset Shape: ',train_set.shape)
print('Test Dataset Shape: ',test_set.shape)
```

```
Train Dataset Shape: (32561, 15)
Test Dataset Shape: (16281, 15)
```

```
[9]: #Checking for NULL values
print('\n***** TRAIN SET *****\n')
print(train_set.info())

print('\n***** TEST SET *****\n')
print(test_set.info())
```

```
***** TRAIN SET *****

<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 32561 entries, 0 to 32560

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	age	32561 non-null	int64
1	workclass	32561 non-null	object
2	fnlwgt	32561 non-null	int64
3	education	32561 non-null	object
4	education_num	32561 non-null	int64
5	marital_status	32561 non-null	object
6	occupation	32561 non-null	object
7	relationship	32561 non-null	object
8	race	32561 non-null	object
9	sex	32561 non-null	object
10	capital_gain	32561 non-null	int64
11	capital_loss	32561 non-null	int64
12	hours_per_week	32561 non-null	int64
13	native_country	32561 non-null	object
14	wage_class	32561 non-null	object

dtypes: int64(6), object(9)

memory usage: 3.7+ MB

None

***** TEST SET *****

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 16281 entries, 0 to 16280

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	age	16281 non-null	int64
1	workclass	16281 non-null	object
2	fnlwgt	16281 non-null	int64
3	education	16281 non-null	object
4	education_num	16281 non-null	int64
5	marital_status	16281 non-null	object
6	occupation	16281 non-null	object
7	relationship	16281 non-null	object
8	race	16281 non-null	object
9	sex	16281 non-null	object
10	capital_gain	16281 non-null	int64
11	capital_loss	16281 non-null	int64
12	hours_per_week	16281 non-null	int64
13	native_country	16281 non-null	object
14	wage_class	16281 non-null	object

dtypes: int64(6), object(9)

memory usage: 1.9+ MB

None

No NULL values are present in train and test dataset, hence no need for imputation techniques.

```
[10]: #Getting an idea of values distribution in train dataset
train_set.describe()
```

```
[10]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	\
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	

	hours_per_week
count	32561.000000
mean	40.437456
std	12.347429
min	1.000000
25%	40.000000
50%	40.000000
75%	45.000000
max	99.000000

```
[11]: #Checking the type of output variable and identifying the type of problem
train_set['wage_class'].value_counts()
```

```
[11]: <=50K    24720
      >50K     7841
      Name: wage_class, dtype: int64
```

Above is a binary classification problem

```
[12]: train_set['workclass'].value_counts()
```

```
[12]: Private          22696
      Self-emp-not-inc  2541
      Local-gov        2093
      ?                1836
      State-gov         1298
      Self-emp-inc       1116
      Federal-gov        960
      Without-pay        14
      Never-worked        7
      Name: workclass, dtype: int64
```

```
[13]: train_set['occupation'].value_counts()
```

```
[13]: Prof-specialty      4140
      Craft-repair      4099
      Exec-managerial   4066
      Adm-clerical      3770
      Sales             3650
      Other-service     3295
      Machine-op-inspct 2002
      ?                1843
      Transport-moving  1597
      Handlers-cleaners 1370
      Farming-fishing   994
      Tech-support      928
      Protective-serv   649
      Priv-house-serv   149
      Armed-Forces      9
      Name: occupation, dtype: int64
```

Seems like few of the missing details are denoted by ?, so replacing the same with `np.nan`

```
[14]: train_set.replace('\?', np.nan, regex=True, inplace=True)
```

```
[15]: train_set['workclass'].value_counts()
```

```
[15]: Private      22696
      Self-emp-not-inc 2541
      Local-gov     2093
      State-gov     1298
      Self-emp-inc  1116
      Federal-gov   960
      Without-pay   14
      Never-worked   7
      Name: workclass, dtype: int64
```

```
[16]: train_set.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             30725 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education_num         32561 non-null  int64
```

```

5 marital_status 32561 non-null object
6 occupation    30718 non-null object
7 relationship   32561 non-null object
8 race          32561 non-null object
9 sex           32561 non-null object
10 capital_gain  32561 non-null int64
11 capital_loss  32561 non-null int64
12 hours_per_week 32561 non-null int64
13 native_country 31978 non-null object
14 wage_class    32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
[17]: #Three columns workclass,occupation,native_country seems to have NULL values.
↳ Lets get the percentage of NaN values in data.
```

```
[18]: nan_cnt=train_set[(train_set.workclass.isna())|(train_set.occupation.
↳ isna())|(train_set.native_country.isna())].count().max()
```

```
[19]: per_nan=nan_cnt/train_set.shape[0] * 100
print('Percentage of NaN values in dataset: ',per_nan,'%')
```

Percentage of NaN values in dataset: 7.367709836921471 %

```
[20]: #Null Values dropped
train_set.dropna(inplace=True)
```

```
[21]: train_set.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 30162 entries, 0 to 32560
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age             30162 non-null  int64
1   workclass       30162 non-null  object
2   fnlwgt         30162 non-null  int64
3   education       30162 non-null  object
4   education_num   30162 non-null  int64
5   marital_status  30162 non-null  object
6   occupation      30162 non-null  object
7   relationship    30162 non-null  object
8   race           30162 non-null  object
9   sex            30162 non-null  object
10  capital_gain    30162 non-null  int64
11  capital_loss    30162 non-null  int64
12  hours_per_week  30162 non-null  int64
13  native_country  30162 non-null  object

```

```
14 wage_class      30162 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

Working with categorical and non numeric data

```
[22]: #Sex Feature
```

```
[23]: train_set.sex.value_counts()
```

```
[23]: Male      20380
      Female    9782
      Name: sex, dtype: int64
```

```
[24]: train_set['sex']=train_set['sex'].str.strip()
      #label encoding
      train_set['sex']=train_set['sex'].map({'Male':1,'Female':0})
```

```
[25]: train_set.sex.value_counts()
```

```
[25]: 1      20380
      0      9782
      Name: sex, dtype: int64
```

```
[26]: #Feature: workclass
```

```
[27]: train_set['workclass']=train_set['workclass'].str.strip()
      train_set.workclass.value_counts()
```

```
[27]: Private      22286
      Self-emp-not-inc  2499
      Local-gov      2067
      State-gov      1279
      Self-emp-inc    1074
      Federal-gov     943
      Without-pay     14
      Name: workclass, dtype: int64
```

```
[28]: df_workclass=pd.
      ↪get_dummies(train_set['workclass'],drop_first=True,prefix='workclass',prefix_sep='_')
```

```
[29]: df_workclass.head()
```

```
[29]:   workclass_Local-gov  workclass_Private  workclass_Self-emp-inc  \
0                    0                    0                    0
1                    0                    0                    0
2                    0                    1                    0
3                    0                    1                    0
```

	0	1	0
workclass_Self-emp-not-inc			
workclass_State-gov			
workclass_Without-pay			
0	0	1	0
1	1	0	0
2	0	0	0
3	0	0	0
4	0	0	0

```
[30]: #Feature: marital_status
```

```
[31]: train_set['marital_status']=train_set['marital_status'].str.strip()
train_set['marital_status'].value_counts()
```

```
[31]: Married-civ-spouse      14065
Never-married              9726
Divorced                   4214
Separated                   939
Widowed                    827
Married-spouse-absent      370
Married-AF-spouse          21
Name: marital_status, dtype: int64
```

```
[32]: df_maritalst=pd.get_dummies(train_set['marital_status'],drop_first=True)
```

```
[33]: df_maritalst.head()
```

```
[33]: Married-AF-spouse  Married-civ-spouse  Married-spouse-absent  \
0                      0                      0                      0
1                      0                      1                      0
2                      0                      0                      0
3                      0                      1                      0
4                      0                      1                      0

Never-married  Separated  Widowed
0              1         0        0
1              0         0        0
2              0         0        0
3              0         0        0
4              0         0        0
```

```
[34]: #Feature: relationship
```

```
[35]: train_set['relationship']=train_set['relationship'].str.strip()
train_set['relationship'].value_counts()
```



```
[35]: Husband          12463
      Not-in-family    7726
      Own-child        4466
      Unmarried        3212
      Wife             1406
      Other-relative    889
      Name: relationship, dtype: int64
```

```
[36]: df_rel=pd.get_dummies(train_set['relationship'],drop_first=True)
```

```
[37]: df_rel.head()
```

```
[37]:   Not-in-family  Other-relative  Own-child  Unmarried  Wife
0              1              0          0          0        0
1              0              0          0          0        0
2              1              0          0          0        0
3              0              0          0          0        0
4              0              0          0          0        1
```

```
[38]: #Feature: occupation
      train_set['occupation']=train_set['occupation'].str.strip()
      train_set['occupation'].value_counts()
```

```
[38]: Prof-specialty      4038
      Craft-repair       4030
      Exec-managerial    3992
      Adm-clerical       3721
      Sales              3584
      Other-service      3212
      Machine-op-inspct  1966
      Transport-moving    1572
      Handlers-cleaners   1350
      Farming-fishing     989
      Tech-support        912
      Protective-serv     644
      Priv-house-serv     143
      Armed-Forces        9
      Name: occupation, dtype: int64
```

```
[39]: df_occ=pd.get_dummies(train_set['occupation'],drop_first=True)
```

```
[40]: df_occ.head()
```

```
[40]:   Armed-Forces  Craft-repair  Exec-managerial  Farming-fishing  \
0              0              0              0              0
1              0              0              1              0
2              0              0              0              0
```

3	0	0	0	0
4	0	0	0	0

	Handlers-cleaners	Machine-op-inspct	Other-service	Priv-house-serv	\
0	0	0	0	0	
1	0	0	0	0	
2	1	0	0	0	
3	1	0	0	0	
4	0	0	0	0	

	Prof-specialty	Protective-serv	Sales	Tech-support	Transport-moving
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	1	0	0	0	0

```
[41]: #Feature: race
train_set['race']=train_set['race'].str.strip()
train_set['race'].value_counts()
```

```
[41]: White                25933
Black                   2817
Asian-Pac-Islander      895
Amer-Indian-Eskimo     286
Other                   231
Name: race, dtype: int64
```

```
[42]: df_race=pd.
↳get_dummies(train_set['race'],drop_first=True,prefix='race',prefix_sep='_')
```

```
[43]: df_race.head()
```


```
[43]:   race_Asian-Pac-Islander  race_Black  race_Other  race_White
0                        0           0           0           1
1                        0           0           0           1
2                        0           0           0           1
3                        0           1           0           0
4                        0           1           0           0
```

```
[44]: #Feature: native_country
```

```
[46]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
[47]: plt.figure(figsize=(20,10))
x=sns.countplot(x=train_set['native_country'])
```

```
x.set_xticklabels(labels=x.get_xticklabels(),rotation=45)
plt.show()
```



[48]: *#Seems like native_country has distribution of US much higher than other hence*
↪dividing it into US and Non-US and label encoding the same.

```
[49]: train_set['native_country']=train_set['native_country'].str.strip()
train_set['native_country']=train_set['native_country'].apply(lambda x:x
↪if('United-States' in x) else 'Non-US')
```

```
[50]: train_set['native_country'].value_counts()
```

```
[50]: United-States    27504
Non-US              2658
Name: native_country, dtype: int64
```

```
[51]: train_set['native_country']=train_set['native_country'].map({'United-States':
↪1, 'Non-US':0})
```

```
[52]: train_set['native_country'].value_counts()
```

```
[52]: 1    27504
0     2658
Name: native_country, dtype: int64
```

```
[53]: #Feature: wage_class
```

```
[54]: train_set['wage_class'].value_counts()
```

```
[54]: <=50K    22654
      >50K     7508
      Name: wage_class, dtype: int64
```

```
[55]: train_set['wage_class']=train_set['wage_class'].str.strip()
```

```
[56]: train_set['wage_class']=train_set['wage_class'].map({'<=50K':1,'>50K':0})
```

```
[57]: train_set['wage_class'].value_counts()
```

```
[57]: 1    22654
      0     7508
      Name: wage_class, dtype: int64
```

```
[58]: train_set.head()
```

```
[58]:   age      workclass  fnlwgt  education  education_num  \
0   39      State-gov   77516    Bachelors             13
1   50  Self-emp-not-inc   83311    Bachelors             13
2   38        Private  215646    HS-grad              9
3   53        Private  234721      11th              7
4   28        Private  338409    Bachelors             13

      marital_status      occupation  relationship   race  sex  \
0      Never-married      Adm-clerical  Not-in-family  White    1
1  Married-civ-spouse  Exec-managerial      Husband  White    1
2        Divorced  Handlers-cleaners  Not-in-family  White    1
3  Married-civ-spouse  Handlers-cleaners      Husband  Black    1
4  Married-civ-spouse  Prof-specialty      Wife  Black    0

      capital_gain  capital_loss  hours_per_week  native_country  wage_class
0           2174             0             40             1             1
1              0             0             13             1             1
2              0             0             40             1             1
3              0             0             40             1             1
4              0             0             40             0             1
```

```
[59]: #Appending all new columns
      new_df=pd.
      ↪concat([train_set,df_workclass,df_maritalst,df_rel,df_occ,df_race],axis=1)
```

```
[60]: new_df.columns
```

```
[60]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
      'marital_status', 'occupation', 'relationship', 'race', 'sex',
```

```

'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
'wage_class', 'workclass_Local-gov', 'workclass_Private',
'workclass_Self-emp-inc', 'workclass_Self-emp-not-inc',
'workclass_State-gov', 'workclass_Without-pay', 'Married-AF-spouse',
'Married-civ-spouse', 'Married-spouse-absent', 'Never-married',
'Separated', 'Widowed', 'Not-in-family', 'Other-relative', 'Own-child',
'Unmarried', 'Wife', 'Armed-Forces', 'Craft-repair', 'Exec-managerial',
'Farming-fishing', 'Handlers-cleaners', 'Machine-op-inspct',
'Other-service', 'Priv-house-serv', 'Prof-specialty', 'Protective-serv',
'Sales', 'Tech-support', 'Transport-moving', 'race_Asian-Pac-Islander',
'race_Black', 'race_Other', 'race_White'],
dtype='object')

```

```

[61]: new_df=new_df.
      ↪drop(columns=['workclass','education','marital_status','occupation','relationship','race'])

```

```

[62]: print(new_df.columns)
      new_df.head()

```

```

Index(['age', 'fnlwgt', 'education_num', 'sex', 'capital_gain', 'capital_loss',
'hours_per_week', 'native_country', 'wage_class', 'workclass_Local-gov',
'workclass_Private', 'workclass_Self-emp-inc',
'workclass_Self-emp-not-inc', 'workclass_State-gov',
'workclass_Without-pay', 'Married-AF-spouse', 'Married-civ-spouse',
'Married-spouse-absent', 'Never-married', 'Separated', 'Widowed',
'Not-in-family', 'Other-relative', 'Own-child', 'Unmarried', 'Wife',
'Armed-Forces', 'Craft-repair', 'Exec-managerial', 'Farming-fishing',
'Handlers-cleaners', 'Machine-op-inspct', 'Other-service',
'Priv-house-serv', 'Prof-specialty', 'Protective-serv', 'Sales',
'Tech-support', 'Transport-moving', 'race_Asian-Pac-Islander',
'race_Black', 'race_Other', 'race_White'],
dtype='object')

```

```

[62]:
  age  fnlwgt  education_num  sex  capital_gain  capital_loss  \
0   39   77516             13    1           2174             0
1   50   83311             13    1              0             0
2   38  215646              9    1              0             0
3   53  234721              7    1              0             0
4   28  338409             13    0              0             0

  hours_per_week  native_country  wage_class  workclass_Local-gov  ...  \
0              40                1           1                  0  ...
1              13                1           1                  0  ...
2              40                1           1                  0  ...
3              40                1           1                  0  ...
4              40                0           1                  0  ...

```

	Priv-house-serv	Prof-specialty	Protective-serv	Sales	Tech-support	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	1	0	0	0	

	Transport-moving	race_Asian-Pac-Islander	race_Black	race_Other	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	1	0	
4	0	0	1	0	

	race_White
0	1
1	1
2	1
3	0
4	0

[5 rows x 43 columns]

All features are now converted to numerical once

```
[96]: test_set.replace('\?', np.nan, regex=True, inplace=True)
test_set.dropna(inplace=True)

test_set['workclass']=test_set['workclass'].str.strip()
test_set['marital_status']=test_set['marital_status'].str.strip()
test_set['relationship']=test_set['relationship'].str.strip()
test_set['occupation']=test_set['occupation'].str.strip()
test_set['race']=test_set['race'].str.strip()

test_set['sex']=test_set['sex'].str.strip()
#label encoding
test_set['sex']=test_set['sex'].map({'Male':1, 'Female':0})

df_workclass=pd.
    ↳get_dummies(test_set['workclass'],drop_first=True,prefix='workclass',prefix_sep='_')

df_maritalst=pd.get_dummies(test_set['marital_status'],drop_first=True)
df_rel=pd.get_dummies(test_set['relationship'],drop_first=True)
df_occ=pd.get_dummies(test_set['occupation'],drop_first=True)
df_race=pd.
    ↳get_dummies(test_set['race'],drop_first=True,prefix='race',prefix_sep='_')
test_set['native_country']=test_set['native_country'].str.strip()
```

```

test_set['native_country']=test_set['native_country'].apply(lambda x:x
↳if('United-States' in x) else 'Non-US')
test_set['native_country']=test_set['native_country'].map({'United-States':
↳1, 'Non-US':0})
test_set['wage_class']=test_set['wage_class'].str.strip()
print(test_set.head())
test_set['wage_class']=test_set['wage_class'].map({'<=50K.':1, '>50K.':0})
test_df=pd.
↳concat([test_set,df_workclass,df_maritalst,df_rel,df_occ,df_race],axis=1)

test_df=test_df.
↳drop(columns=['workclass','education','marital_status','occupation','relationship','race'])

```

	age	workclass	fnlwgt	education	education_num	marital_status	\
0	25	Private	226802	11th	7	Never-married	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	
3	44	Private	160323	Some-college	10	Married-civ-spouse	
5	34	Private	198693	10th	6	Never-married	

	occupation	relationship	race	sex	capital_gain	capital_loss	\
0	Machine-op-inspct	Own-child	Black	1	0	0	
1	Farming-fishing	Husband	White	1	0	0	
2	Protective-serv	Husband	White	1	0	0	
3	Machine-op-inspct	Husband	Black	1	7688	0	
5	Other-service	Not-in-family	White	1	0	0	

	hours_per_week	native_country	wage_class
0	40	1	<=50K.
1	50	1	<=50K.
2	40	1	>50K.
3	40	1	>50K.
5	30	1	<=50K.

[97]: test_df.columns

[97]: Index(['age', 'fnlwgt', 'education_num', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'native_country', 'wage_class', 'workclass_Local-gov', 'workclass_Private', 'workclass_Self-emp-inc', 'workclass_Self-emp-not-inc', 'workclass_State-gov', 'workclass_Without-pay', 'Married-AF-spouse', 'Married-civ-spouse', 'Married-spouse-absent', 'Never-married', 'Separated', 'Widowed', 'Not-in-family', 'Other-relative', 'Own-child', 'Unmarried', 'Wife', 'Armed-Forces', 'Craft-repair', 'Exec-managerial', 'Farming-fishing', 'Handlers-cleaners', 'Machine-op-inspct', 'Other-service', 'Priv-house-serv', 'Prof-specialty', 'Protective-serv', 'Sales', 'Tech-support', 'Transport-moving', 'race_Asian-Pac-Islander',

```

        'race_Black', 'race_Other', 'race_White'],
        dtype='object')

```

```
[98]: test_df.head()
```

```

[98]:   age  fnlwgt  education_num  sex  capital_gain  capital_loss  \
0    25  226802             7    1           0           0
1    38   89814             9    1           0           0
2    28  336951            12    1           0           0
3    44  160323            10    1          7688           0
5    34  198693             6    1           0           0

      hours_per_week  native_country  wage_class  workclass_Local-gov  ...  \
0                40                1           1                   0  ...
1                50                1           1                   0  ...
2                40                1           0                   1  ...
3                40                1           0                   0  ...
5                30                1           1                   0  ...

      Priv-house-serv  Prof-specialty  Protective-serv  Sales  Tech-support  \
0                   0                0                0     0             0
1                   0                0                0     0             0
2                   0                0                1     0             0
3                   0                0                0     0             0
5                   0                0                0     0             0

      Transport-moving  race_Asian-Pac-Islander  race_Black  race_Other  \
0                   0                        0           1           0
1                   0                        0           0           0
2                   0                        0           0           0
3                   0                        0           1           0
5                   0                        0           0           0

      race_White
0              0
1              1
2              1
3              0
5              1

```

```
[5 rows x 43 columns]
```

```
[65]: new_df.columns
```

```

[65]: Index(['age', 'fnlwgt', 'education_num', 'sex', 'capital_gain', 'capital_loss',
            'hours_per_week', 'native_country', 'wage_class', 'workclass_Local-gov',
            'workclass_Private', 'workclass_Self-emp-inc',

```



```

'workclass_Self-emp-not-inc', 'workclass_State-gov',
'workclass_Without-pay', 'Married-AF-spouse', 'Married-civ-spouse',
'Married-spouse-absent', 'Never-married', 'Separated', 'Widowed',
'Not-in-family', 'Other-relative', 'Own-child', 'Unmarried', 'Wife',
'Armed-Forces', 'Craft-repair', 'Exec-managerial', 'Farming-fishing',
'Handlers-cleaners', 'Machine-op-inspct', 'Other-service',
'Priv-house-serv', 'Prof-specialty', 'Protective-serv', 'Sales',
'Tech-support', 'Transport-moving', 'race_Asian-Pac-Islander',
'race_Black', 'race_Other', 'race_White'],
dtype='object')

```

```

[99]: train_x=new_df.drop(columns=['wage_class'])
      train_y=new_df['wage_class']
      test_x=test_df.drop(columns=['wage_class'])
      test_y=test_df['wage_class']

```

```

[100]: #Applying models on machine learning

```

```

[101]: from sklearn.model_selection import train_test_split,GridSearchCV
      from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, \
      ↪roc_auc_score
      import xgboost as xgb

```

```

[102]: #Fitting model on training data
      model = xgb.XGBClassifier(objective='binary:logistic')
      model.fit(train_x, train_y)

```

```

[102]: XGBClassifier()

```

```

[103]: # cheking training accuracy
      y_pred = model.predict(train_x)
      predictions = [round(value) for value in y_pred]
      accuracy = accuracy_score(train_y,predictions)
      accuracy

```

```

[103]: 0.8634374378356873

```

```

[104]: # checking initial test accuracy
      y_pred = model.predict(test_x)
      predictions = [round(value) for value in y_pred]
      accuracy = accuracy_score(test_y,predictions)
      accuracy

```

```

[104]: 0.8638778220451527

```

Model seems not be overfitting as test and train data accuracy is almost same.

```
[105]: from sklearn.model_selection import GridSearchCV
```

```
[106]: param_grid={  
    ' learning_rate': [1,0.5,0.1,0.01,0.001],  
    'max_depth': [3,5,10,20],  
    'n_estimators': [10,50,100,200]  
}
```

```
[110]: grid= GridSearchCV(xgb.XGBClassifier(objective='binary:logistic'),param_grid,verbose=3, n_jobs=-1)
```

```
[111]: grid.fit(train_x,train_y)
```

Fitting 5 folds for each of 80 candidates, totalling 400 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 16 tasks | elapsed: 15.7s

[Parallel(n_jobs=-1)]: Done 112 tasks | elapsed: 3.6min

[Parallel(n_jobs=-1)]: Done 272 tasks | elapsed: 8.5min

[Parallel(n_jobs=-1)]: Done 400 out of 400 | elapsed: 13.2min finished

```
[111]: GridSearchCV(estimator=XGBClassifier(), n_jobs=-1,  
    param_grid={' learning_rate': [1, 0.5, 0.1, 0.01, 0.001],  
    'max_depth': [3, 5, 10, 20],  
    'n_estimators': [10, 50, 100, 200]},  
    verbose=3)
```

```
[112]: # To find the parameters giving maximum accuracy  
grid.best_params_
```

```
[112]: {' learning_rate': 1, 'max_depth': 5, 'n_estimators': 200}
```

```
[122]: # Create new model using the same parameters  
new_model=xgb.XGBClassifier(learning_rate= 1, max_depth= 5, n_estimators= 50)  
new_model.fit(train_x, train_y)
```

```
[122]: XGBClassifier(learning_rate=1, max_depth=5, n_estimators=50)
```

```
[123]: y_pred_new = new_model.predict(test_x)  
predictions_new = [round(value) for value in y_pred_new]  
accuracy_new = accuracy_score(test_y,predictions_new)  
accuracy_new
```

```
[123]: 0.8602921646746348
```

```
[124]: # cheking training accuracy  
y_pred = new_model.predict(train_x)  
predictions = [round(value) for value in y_pred]  
accuracy = accuracy_score(train_y,predictions)  
accuracy
```

[124]: 0.9017306544658842

```
[125]: # As we have increased the accuracy of the model, we'll save this model  
import pickle  
filename = 'xgboost_model.pickle'  
pickle.dump(new_model, open(filename, 'wb'))
```